

```
[j111046@login-2-1 PA3]$ hadoop jar pageRank.jar PageRank class3/input.txt class3/output
19/09/23 18:19:46 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
19/09/23 18:19:46 INFO input.FileInputFormat: Total input paths to process : 1
19/09/23 18:19:47 INFO mapreduce.JobSubmitter: number of splits:1
19/09/23 18:19:47 INFO Configuration.deprecation: mapred.textoutputformat.separator is deprecated. Instead, use mapreduce.output.textoutputformat.separator
19/09/23 18:19:47 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1569268517451_0268
19/09/23 18:19:47 INFO impl.YarnClientImpl: Submitted application application_1569268517451_0268
19/09/23 18:19:47 INFO mapreduce.Job: The url to track the job: http://babar.es.its.nyu.edu:8088/proxy/application_1569268517451_0268/
19/09/23 18:19:47 INFO mapreduce.Job: Running job: job_1569268517451_0268
19/09/23 18:19:52 INFO mapreduce.Job: Job job_1569268517451_0268 running in uber mode : false
19/09/23 18:19:52 INFO mapreduce.Job: map 0% reduce 0%
19/09/23 18:19:56 INFO mapreduce.Job: map 100% reduce 0%
19/09/23 18:20:01 INFO mapreduce.Job: map 100% reduce 100%
19/09/23 18:20:02 INFO mapreduce.Job: Job job_1569268517451_0268 completed successfully
19/09/23 18:20:02 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=188
    FILE: Number of bytes written=309041
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=207
    HDFS: Number of bytes written=151
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=8360
    Total time spent by all reduces in occupied slots (ms)=14070
    Total time spent by all map tasks (ms)=2090
    Total time spent by all reduce tasks (ms)=2345
    Total vcore-milliseconds taken by all map tasks=2090
    Total vcore-milliseconds taken by all reduce tasks=2345
    Total megabyte-milliseconds taken by all map tasks=8560640
    Total megabyte-milliseconds taken by all reduce tasks=14407680
  Map-Reduce Framework
    Map input records=6
    Map output records=21
    Map output bytes=257
    Map output materialized bytes=184
    Input split bytes=107
    Combine input records=0
    Combine output records=0
```

```

        Reduce input groups=6
        Reduce shuffle bytes=184
        Reduce input records=21
        Reduce output records=6
        Spilled Records=42
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=95
        CPU time spent (ms)=1910
        Physical memory (bytes) snapshot=911691776
        Virtual memory (bytes) snapshot=7472320512
        Total committed heap usage (bytes)=2375024640
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Input Format Counters
        Bytes Read=100
    File Output Format Counters
        Bytes Written=151
[jl11046@login-2-1 PA3]$ hdfs dfs -ls class3/output
Found 2 items
-rw-r--r--+ 3 jl11046 users      0 2019-09-23 18:20 class3/output/_SUCCESS
-rw-r--r--+ 3 jl11046 users    151 2019-09-23 18:19 class3/output/part-r-00000
[jl11046@login-2-1 PA3]$ hdfs dfs -cat class3/part-r-00000
cat: `class3/part-r-00000': No such file or directory
[jl11046@login-2-1 PA3]$ hdfs dfs -cat class3/output/part-r-00000
A C J 0.116669
B D E J 0.20000040000000002
C A B 0.20000040000000002
D A B C E J 0.05555566666666667
E J 0.08888906666666667
J B C 0.33888956666666667

```

```
[jll11046@login-2-1 PA3]$ hadoop jar pageRank.jar PageRank class3/output/part-r-00000 class3/output2
19/09/23 18:20:49 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
19/09/23 18:20:50 INFO input.FileInputFormat: Total input paths to process : 1
19/09/23 18:20:50 INFO mapreduce.JobSubmitter: number of splits:1
19/09/23 18:20:50 INFO Configuration.deprecation: mapred.textoutputformat.separator is deprecated. Instead, use mapreduce.output.textoutputformat.separator
19/09/23 18:20:50 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1569268517451_0270
19/09/23 18:20:51 INFO impl.YarnClientImpl: Submitted application application_1569268517451_0270
19/09/23 18:20:51 INFO mapreduce.Job: The url to track the job: http://bahar.es.its.nyu.edu:8088/proxy/application_1569268517451_0270/
19/09/23 18:20:51 INFO mapreduce.Job: Running job: job_1569268517451_0270
19/09/23 18:20:55 INFO mapreduce.Job: Job job_1569268517451_0270 running in uber mode : false
19/09/23 18:20:55 INFO mapreduce.Job: map 0% reduce 0%
19/09/23 18:21:00 INFO mapreduce.Job: map 100% reduce 0%
19/09/23 18:21:05 INFO mapreduce.Job: map 100% reduce 100%
19/09/23 18:21:06 INFO mapreduce.Job: Job job_1569268517451_0270 completed successfully
19/09/23 18:21:06 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=243
    FILE: Number of bytes written=309173
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=268
    HDFS: Number of bytes written=161
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=8624
    Total time spent by all reduces in occupied slots (ms)=14064
    Total time spent by all map tasks (ms)=2156
    Total time spent by all reduce tasks (ms)=2344
    Total vcore-milliseconds taken by all map tasks=2156
    Total vcore-milliseconds taken by all reduce tasks=2344
    Total megabyte-milliseconds taken by all map tasks=8830976
    Total megabyte-milliseconds taken by all reduce tasks=14401536
  Map-Reduce Framework
    Map input records=6
    Map output records=21
    Map output bytes=365
    Map output materialized bytes=239
    Input split bytes=117
```

```

Combine input records=0
Combine output records=0
Reduce input groups=6
Reduce shuffle bytes=239
Reduce input records=21
Reduce output records=6
Spilled Records=42
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=94
CPU time spent (ms)=1920
Physical memory (bytes) snapshot=911024128
Virtual memory (bytes) snapshot=7473655808
Total committed heap usage (bytes)=2416443392

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=151
File Output Format Counters
  Bytes Written=161
[jl11046@login-2-1 PA3]$ hdfs dfs -ls class3/output2
Found 2 items
-rw-r--r--+ 3 jl11046 users          0 2019-09-23 18:21 class3/output2/_SUCCESS
-rw-r--r--+ 3 jl11046 users        161 2019-09-23 18:21 class3/output2/part-r-00000
[jl11046@login-2-1 PA3]$ hdfs dfs -cat class3/output2/part-r-00000
A C J 0.11111133333333334
B D E J 0.2805561166666667
C A B 0.23888936666666669
D A B C E J 0.06666680000000001
E J 0.07777793333333334
J B C 0.22500045000000002

```

```
[j111046@login-2-1 PA3]$ hadoop jar pageRank.jar PageRank class3/output2/part-r-00000 class3/output3
19/09/23 18:21:33 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
19/09/23 18:21:34 INFO input.FileInputFormat: Total input paths to process : 1
19/09/23 18:21:34 INFO mapreduce.JobSubmitter: number of splits:1
19/09/23 18:21:34 INFO Configuration.deprecation: mapred.textoutputformat.separator is deprecated. Instead, use mapreduce.output.textoutputformat.separator
19/09/23 18:21:34 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1569268517451_0272
19/09/23 18:21:35 INFO impl.YarnClientImpl: Submitted application application_1569268517451_0272
19/09/23 18:21:35 INFO mapreduce.Job: The url to track the job: http://babar.es.its.nyu.edu:8088/proxy/application_1569268517451_0272/
19/09/23 18:21:35 INFO mapreduce.Job: Running job: job_1569268517451_0272
19/09/23 18:21:40 INFO mapreduce.Job: Job job_1569268517451_0272 running in uber mode : false
19/09/23 18:21:40 INFO mapreduce.Job: map 0% reduce 0%
19/09/23 18:21:44 INFO mapreduce.Job: map 100% reduce 0%
19/09/23 18:21:49 INFO mapreduce.Job: map 100% reduce 100%
19/09/23 18:21:50 INFO mapreduce.Job: Job job_1569268517451_0272 completed successfully
19/09/23 18:21:50 INFO mapreduce.Job: Counters: 49
    File System Counters
      FILE: Number of bytes read=250
      FILE: Number of bytes written=309189
      FILE: Number of read operations=0
      FILE: Number of large read operations=0
      FILE: Number of write operations=0
      HDFS: Number of bytes read=279
      HDFS: Number of bytes written=160
      HDFS: Number of read operations=6
      HDFS: Number of large read operations=0
      HDFS: Number of write operations=2
    Job Counters
      Launched map tasks=1
      Launched reduce tasks=1
      Data-local map tasks=1
      Total time spent by all maps in occupied slots (ms)=8620
      Total time spent by all reduces in occupied slots (ms)=13590
      Total time spent by all map tasks (ms)=2155
      Total time spent by all reduce tasks (ms)=2265
      Total vcore-milliseconds taken by all map tasks=2155
      Total vcore-milliseconds taken by all reduce tasks=2265
      Total megabyte-milliseconds taken by all map tasks=8826880
      Total megabyte-milliseconds taken by all reduce tasks=13916160
    Map-Reduce Framework
      Map input records=6
      Map output records=21
      Map output bytes=383
      Map output materialized bytes=246
      Input split bytes=118
      Combine input records=0
```

```

Combine output records=0
Reduce input groups=6
Reduce shuffle bytes=246
Reduce input records=21
Reduce output records=6
Spilled Records=42
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=88
CPU time spent (ms)=1780
Physical memory (bytes) snapshot=910516224
Virtual memory (bytes) snapshot=7470661632
Total committed heap usage (bytes)=2390228992

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=161

File Output Format Counters
  Bytes Written=160

[jl11046@login-2-1 PA3]$ hdfs dfs -ls class3/output3
Found 2 items
-rw-r--r--+ 3 jl11046 users          0 2019-09-23 18:21 class3/output3/_SUCCESS
-rw-r--r--+ 3 jl11046 users        160 2019-09-23 18:21 class3/output3/part-r-00000
[jl11046@login-2-1 PA3]$ hdfs dfs -cat class3/output3/part-r-00000
A C J 0.13277804333333335
B D E J 0.24527826833333338
C A B 0.18138925166666667
D A B C E J 0.09351870555555557
E J 0.10685206555555557
J B C 0.24018566555555556
[jl11046@login-2-1 PA3]$

```